

The Undrawn Book

Applying Basel III credit-card rules to the quarterly regulatory filings of 50 small US banks — and why most of the exposure sits where nobody currently counts it.

3,682 Basel paragraphs parsed 673 rules extracted, all grounded 600 bank-quarters computed
2023 Q1 – 2025 Q4

The brief asks for four things in order: extract the rules, classify them, derive parameters from them, then find those parameters in banks' quarterly reports. This document covers that chain as it now stands, end to end, with a worked example. Everything described here runs — the figures below are produced by the code, not estimated.

01 – Inputs

Two corpora, and the join between them

The project needs a regulatory side and a bank side, and the difficulty the supervisor identified is real: the quarterly report is structured, but it is not written in the vocabulary of the rules.

Regulatory side

- BIS consolidated Basel Framework, 1,982 pages, parsed to 3,682 paragraphs, each keeping its real citation (CRE20.68) and heading path.
- 673 rules extracted across 23 chapters — 367 with a number in them, 144 bearing on cards.
- Classified by kind: capital, liquidity, credit-risk RWA, concentration, provisioning, issuance, disclosure.

Bank side

- FFIEC Call Reports — the mandatory quarterly filing — pulled from the CDR bulk distribution.
- 50 banks × 12 quarters = 600 filings, complete panel, no gaps.
- Small banks: under \$10bn in assets, the community bank leverage-ratio threshold. Median \$4.1bn.
- Each filing is ~1,750 fields keyed by MDRM code (RCON3815).

The join

A rule cites a Basel paragraph; a filing reports an MDRM code. The bridge is a fixed vocabulary of named quantities, each bound to specific codes — so a formula written against a rule is executable against a filing. A formula naming `EAD_transactor` cannot be run against anything; one naming `credit_card_balances` resolves to `RCFDB538` or `RCONB538` and produces a number.

02 — Method

Five steps, each resting on a cited rule

The steps are a genuine sequence: each consumes the previous one's output. The rail on the left of each step names the provisions it rests on.

1

CRE20.94

CRE20.100

12 CFR 217.33(b)(1)

12 CFR 217.2

Exposure at default

Basel converts off-balance-sheet commitments into credit-equivalent exposure with a credit conversion factor. Under the Basel III finalisation standard, effective 1 January 2023, a commitment that is *unconditionally cancellable at any time without prior notice* takes a **10% CCF**.

$$\text{EAD} = \text{drawn} + \text{undrawn} \times \text{CCF}$$

Comenity: $\$6.76\text{bn} + \$37.77\text{bn} \times 0.10 = \mathbf{\$10.53\text{bn}}$

Two qualifications the framework imposes on itself. Basel does not classify card lines as unconditionally cancellable — that turns on the contract and on local law, and CRE20.100 expressly lets supervisors impose a *higher* factor where cancellation is constrained in practice. In the US the classification does generally hold: Regulation Z requires no advance notice to cut a card line, and issuers report substantially all unused card lines as unconditionally cancelable in Schedule RC-R.

The US rule in force applies 0% to the same lines (`12 CFR 217.33(b)(1)`), so they contribute nothing to reported exposure. The 10% figure exists in the US only as a March 2026 proposal covering the largest banks; the companion proposal for everyone else retains 0%.

2

CRE20.65

CRE20.66

CRE20.68

CRE20.106

12 CFR 217.32(1)(5)

Risk weight

Credit cards are named explicitly in Basel's "regulatory retail" product criterion, so the book qualifies subject to a €1m per-counterparty ceiling and a 0.2% granularity limit. The weight then turns on whether the obligor is a *transactor* — someone who repaid in full at every scheduled date for twelve months.

45% transactor · **75%** regulatory retail, non-transactor · **100%** other retail

No US regulatory report discloses the transactor split, so the output is a **band across all three weights** rather than a point estimate. 75% is the default rather than a conservative election — it is where regulatory retail lands when transactor status cannot be evidenced — and it is not an upper bound either: defaulted exposures attract 150% under CRE20.106.

The US assigns credit-card loans to no retail category at all. They fall into the residual "other assets" bucket at **100%** (12 CFR 217.32(1)(5)), rising to **150%** on the unsecured portion of balances 90+ days past due or on nonaccrual (§ 217.32(k)).

3

RBC20.1

RBC30.2

RC-R P859

RC-R A223

Capital adequacy

The card book's Basel RWA replaces its US contribution in the bank's reported total, and the restated ratio is tested against the minimum and the buffer.

$$RWA' = RWA_{\text{reported}} - (\text{drawn} \times 100\%) + (\text{EAD} \times rw)$$

test **CET1** $\geq$ **4.5%** (minimum) and $\geq$ **7.0%** (with the 2.5% conservation buffer)

Falling into the buffer range constrains capital distributions rather than breaching the minimum, so the two thresholds are reported separately. 7.0% is the effective CET1 threshold only for a bank meeting its tier 1 and total capital minimums entirely with other instruments; CET1 must first fill any shortfall against those, and the US buffer is the lowest of the three legs (12 CFR 217.11(a)(3)(i)).

4

LCR40.59

LCR40.64

LCR40.67

RC 0071 / 0081

Liquidity stress

Basel assigns a 5% **drawdown** within thirty days to undrawn committed facilities to retail and small-business customers. But **LCR40.59 excludes unconditionally cancellable facilities from that section altogether**, routing them to national discretion under LCR40.67 — the same property that earns a card line its low conversion factor also removes it from the committed-facility drawdown.

So the 5% here is a **chosen stress parameter, not a Basel requirement**: it borrows the rate Basel sets for facilities a bank cannot cancel and applies it to facilities it can. That makes it deliberately severe, and makes the rate itself something to defend — which is exactly the scenario-plausibility question the brief raises.

$$\text{cover} = \text{liquid assets} \div (\text{undrawn} \times 0.05)$$

Comenity: \$1.23bn ÷ \$1.89bn = **0.65×**

It is not a compliance test either way. The US LCR sets scope by tailoring category rather than a single asset test, and no applicability hook reaches an institution below \$10bn.

5

RC-N B575 / B576

RI-B B514 / B515

RI B485

RC-K B561

Credit quality and economics

Delinquency, loss rate, portfolio yield and allowance coverage — computed against the *average* balance rather than period-end, since income items are flows.

Comenity: 30–89 days **3.4%** · 90+ days **3.6%** · net charge-offs **7.0%** · yield **26.8%** · allowance/drawn **11.4%**

These are the inputs to the performance criterion the brief asks us to define, and they are what a redistribution proposal has to preserve.

03 — Worked example

Comenity Bank, year-end 2025

A \$7.8bn Delaware bank whose loan book is 100% credit cards — and whose committed but undrawn lines are five times its balance sheet.

Same bank, same filing, four regulatory treatments. All figures computed from the 2025-12-31 Call Report.

TREATMENT	EAD	CARD RWA	TOTAL RWA	CET1 RATIO	STATUS
US rule in force today	\$6.76bn	\$6.76bn	\$6.98bn	15.1%	above buffer
Basel @ 45% — all transactors	\$10.53bn	\$4.74bn	\$4.96bn	21.2%	above buffer
Basel @ 75% — central case	\$10.53bn	\$7.90bn	\$8.12bn	13.0%	above buffer
Basel @ 100% — other retail	\$10.53bn	\$10.53bn	\$10.76bn	9.8%	above buffer

A separate figure answers a different question: converting the undrawn lines alone, while leaving the drawn book at its current US treatment, gives **10.7%**. That isolates the exposure nobody counts today. The 13.0% row is what Basel would actually require, because Basel also moves the drawn balances from 100% down to 75%.

What the example demonstrates

The bank is comfortably capitalised on every measure — and its liquidity cover on its own committed lines is **0.65×**. The binding constraint is not capital. It is that a 5% drawdown of promises already made would exhaust the cash and securities on hand. That is the question the supervisor posed, answered with the bank's own filing.

04 — Coverage

Across all 50 banks

Selected results at 2025-12-31. Ordinary community banks barely move — their card books are small. The specialists move enormously.

BANK	ASSETS	UNDRAWN LINES	CET1 REPORTED	BASEL @ 75%	LIQUIDITY COVER
Credit First NA	\$0.04bn	\$11.44bn	387.7%	4.3%	0.06×
Comenity Bank	\$7.76bn	\$37.77bn	15.1%	13.0%	0.65×
John Deere Financial	\$3.40bn	\$33.50bn	—	CBLR	0.00×
Credit One Bank	\$1.97bn	\$8.11bn	—	CBLR	0.53×
Nicolet National Bank	\$9.17bn	\$0.14bn	11.2%	11.1%	—
Peoples Bank	\$9.64bn	\$0.09bn	12.0%	12.0%	—

6 / 49

banks cannot cover a 5% drawdown of their own committed card lines from cash and securities

11.9 → 14.3%

line utilisation across the cohort, 2023 to 2025, while committed capacity stayed flat at ~\$130bn

05 – Limits

What is not computable, and why that matters

The most useful output of a data-grounded method is an honest account of where it stops. Three limits are structural, not fixable by better engineering, and each is a finding in its own right.

Availability of the quantities Basel's card rules require, across the 50 banks at 2025-12-31.

QUANTITY	SOURCE	COVERAGE	CONSEQUENCE
Unused card lines	RC-L 3815	49 / 50	usable
Card balances	RC-C B538	45 / 50	usable
Card interest income	RI B485	43 / 50	usable
Delinquency 30–89 / 90+	RC-N B575/B576	37 / 26	partial
Risk-weighted assets	RC-R A223	42 / 50	8 CBLR electors report none
Transactor / revolver split	—	0 / 50	not collected anywhere
Card book by PD band	RC-O N040–N054	0 / 50	withheld from all 4,394 filers
HQLA, net 30-day outflows	FR 2052a	0 / 50	\$100bn+ filers only

Why each limit is itself a result

The transactor split decides Basel's headline retail weight and no one collects it. A rule that turns on an unobservable is unenforceable from outside the issuer — which is why the output is a band.

Eight of fifty banks report no risk-weighted assets at all. Having elected the community bank leverage ratio, they are outside the risk-based framework entirely. Basel's risk weights are not hard to check for them; they do not apply. That is the supervisor's own point — requirements scale with size — showing up as missing data.

The data that would permit a proper risk segmentation exists and is withheld from everyone. Schedule RC-O buckets the card book by two-year probability of default. Those fields return the literal marker `CONF` for all 4,394 filers — the FDIC does not release them institution by institution, at any size. Nothing can be inferred from their absence, and the analysis must use the standardised approach rather than internal ratings.

The community bank leverage ratio is an election, not a size band. Being under \$10bn does not imply a bank reports no risk-weighted assets — the asset test is one of four qualifying criteria and the framework is opt-in. Nationally, 1,722 of 4,394 filers leave the RWA field empty, so a capital panel built naively on it drops 39% of banks, skewed small.

The divergence to settle before the next phase

Basel and the US rule disagree on precisely the provisions this project centres on. Basel splits regulatory retail 45%/75% by transactor status; the US applies a flat 100% and does not recognise the split. Basel converts unconditionally cancellable commitments at 10%; the US converts them at 0% — so the **\$128.6bn** of undrawn card lines across these 50 banks attracts no capital at all under the rules they actually report against.

Applying Basel literally to Call Report data will therefore not reproduce any US bank's published ratios. The question for you: extract Basel as specified, the US implementation the filings reflect, or both with an explicit divergence map? The third is more work and is the only one that makes the comparison defensible.

06 — Assurance

How we know the numbers are right

Rules are grounded mechanically. Every extracted rule carries a quotation that must appear character-for-character in the paragraph it cites, checked by exact substring match. No model is asked whether a quote is faithful, so no model can be wrong about it. All 673 rules pass; those that fail are dropped, not repaired.

Field mappings are cross-validated. The CET1 ratio computed from our own mapping reproduces the ratio each bank reports itself. Where it did not, the mapping was wrong and was fixed.

Regulatory assumptions are tested against the filing, not asserted. Schedule RC-R Part II reports Comenity's undrawn card lines at their full \$37.77bn face amount with a credit-equivalent amount of exactly zero. That is the US 0% conversion factor visible as a reported line item, which is stronger evidence than the ratio of RWA to drawn balances we first used.

Implausible results are chased, not shipped. John Deere Financial appeared to hold zero liquid assets; the cause was a field that not every bank files, and correcting it changed the population failing the liquidity stress from seven banks to six.

What the verification pass changed

Twenty-four claims in this document were checked against the Code of Federal Regulations and the parsed framework before publication. Four were corrected as a result, and they are worth knowing about because each was wrong in a way that would have survived casual review.

The claim that a card line simply *is* an unconditionally cancellable commitment was an inference dressed as a definition — Basel says no such thing, and the US definition omits Basel's "without prior notice" limb. The 5% liquidity drawdown turned out to be excluded from the very lines it was being applied to, which reframes step 4 from a requirement into a chosen stress. The US retail weight comes from the residual "other assets" bucket, not a retail category, and carries a 150% band this document had omitted. And the claim that RC-O's PD bands are blank because small banks are exempt was simply false: they are withheld from every filer at every size.

One thing still moves. US adoption of the Basel III endgame was re-proposed in March 2026 and nothing has been finalised, so every statement here about the US rule describes what is in force as of August 2026, not what is coming.

07 — Next

What this sets up

Steps 1–5 produce a comparable, rule-grounded profile for every bank in every quarter. That is the input the rest of the brief needs.

Performance criterion and ranking. Yield, loss rate, capital headroom and liquidity cover are all computed; what remains is choosing and defending a weighting, then ranking within the size tier.

Benchmark and redistribute. The cohort is already usefully split — Comenity withdrew \$14.2bn of committed lines over three years while Merrick grew its book 93%. Those are opposed strategies at comparable scale, which is better material than a peer group that all behaves alike.

Scenarios. The five steps are the objective function a scenario perturbs: a utilisation shock moves step 1, a migration to "other retail" moves step 2, a funding shock moves step 4. Each is bounded by something observed in the panel, which is the plausibility check the brief requires.